

幻灯片标题

作者姓名

作者机构

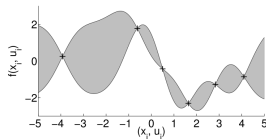
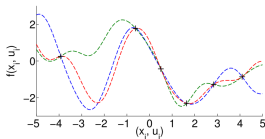
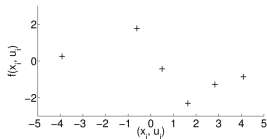
2020/01/01

Outline

- Introduction
- PILCO
- DeepPILCO
- PETS

Motivation

model-based 方法有着较好的 efficiency, 但存在一定的局限性: model bias (对于小样本 & 无先验, 这个问题更为严重), 使用有 bias 的 model 去预测, 进而会带来更大的误差。



Solution

学习一个 probabilistic dynamics model, 这样的 model 能够 express uncertainty。在 PILCO 中, 使用 GP regression 来学习这个 probabilistic dynamics model。

believe 刚才学到的 model, 不做 **rollout**, 而是直接计算

$$J^\pi = \mathbb{E}_{\mathbf{x}_t} [c(\mathbf{x}_t)] = \int c(\mathbf{x}_t) \mathcal{N}(\mathbf{x}_t | \mu_t, \Sigma_t) d\mathbf{x}_t$$

之后采用基于梯度 ($\frac{dJ}{d\theta}$) 的 Policy Search, 找到最优参数, 使得

$$\theta^* = \arg \min_{\theta} J^{\pi_{\theta}}$$

此时得到 $\pi^* = \pi_{\theta^*}$

Algorithm 1 PILCO

- 1: **init:** Sample controller parameters $\theta \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$.
Apply random control signals and record data.
 - 2: **repeat**
 - 3: Learn probabilistic (GP) dynamics model, see Sec. 2.1, using all data.
 - 4: Model-based policy search, see Sec. 2.2–2.3.
 - 5: **repeat**
 - 6: Approximate inference for policy evaluation, see Sec. 2.2: get $J^\pi(\theta)$, Eqs. (10)–(12), (24).
 - 7: Gradient-based policy improvement, see Sec. 2.3: get $dJ^\pi(\theta)/d\theta$, Eqs. (26)–(30).
 - 8: Update parameters θ (e.g., CG or L-BFGS).
 - 9: **until** convergence; **return** θ^*
 - 10: Set $\pi^* \leftarrow \pi(\theta^*)$.
 - 11: Apply π^* to system (single trial/episode) and record data.
 - 12: **until** task learned
-

Improving PILCO with Bayesian neural network dynamics models

Motivation

GP regression 不适合对复杂环境 (e.g. 高维度) 建模。

Solution

主要考虑是将 regression 替换为 NN (Neural Network), 但需要解决带来的两个问题: Output Uncertainty & Input Uncertainty

① Output Uncertainty

将 GP regression 替换为 NN, 能够进行更复杂的建模, 但一般的 NN 不能 express model uncertainty, 考虑使用 BNN (Bayesian Neural Network)。

② Input Uncertainty

PILCO 因为使用了 GP regression, 因此可以 analytically propagates state distributions through the dynamics model, i.e. 推导出分布 $p(X_0), \dots, p(X_T)$, 这些分布能够表示不确定的 input X_0, \dots, X_T 。

但若使用 BNN, 他的输入层只能接受确定性的 input, 无法表示所需要的不确定性, 因此考虑使用 particle methods:

Algorithm 2 Step 6 of Algorithm 1: *Predict* system trajectories from $p(X_0)$ to $p(X_T)$

- 1: *Define* time horizon T .
 - 2: *Initialise* set of K particles $x_0^k \sim P(X_0)$.
 - 3: **for** $k = 1$ to K **do**
 - 4: Sample BNN dynamics model weights W^k .
 - 5: **end for**
 - 6: **for** time $t = 1$ to T **do**
 - 7: **for** each particle x_t^1 to x_t^K **do**
 - 8: Evaluate BNN with weights W^k and input particle x_t^k , obtain output y_t^k .
 - 9: **end for**
 - 10: Calculate mean μ_t and standard deviation σ_t^2 of $\{y_t^1, \dots, y_t^K\}$.
 - 11: Sample set of K particles $x_{t+1}^k \sim \mathcal{N}(\mu_t, \sigma_t^2)$.
 - 12: **end for**
-

Deep Reinforcement Learning in a Handful of Trials using Probabilistic Dynamics Models

Motivation

和上一篇一样在解决 PILCO 的 complexity 问题

Solution

整体的核心思想和上一篇 DeepPILCO 一样：用更复杂的 NN 来代替 GP regression，并解决所带来的 uncertainty 问题。不过这一篇的结果要比上一篇好很多。

1. PE (Probabilistic Ensemble)

🔊 aleatoric (inherent system stochasticity)

使用 probabilistic NN 来解决 aleatoric uncertainty (上一篇的 output uncertainty)

与 Bayesian NN 不同的是, 这篇文章使用了 log prediction probability 来表达 uncertainty:

$$\text{loss}_p(\theta) = - \sum_{n=1}^N \log \tilde{f}_\theta(s_{n+1} \mid s_n, a_n)$$

该 NN 的 input 是 s, a , output 为 \tilde{f}_θ , 是一个 parameterized distribution。

② epistemic (subjective uncertainty, due to limited data)

使用 Ensemble 来解决 epistemic uncertainty

核心思想就是对多个 Probabilistic NN 进行 ensemble: consider ensembles of B -many bootstrap models, using θ_b to refer to the parameters of our b^{th} model \tilde{f}_{θ_b} , then get $\tilde{f}_{\theta} = \frac{1}{B} \sum_{b=1}^B \tilde{f}_{\theta_b}$.

2.TS (trajectory sampling)

create P particles from the current state, $s_{t=0}^p = s_0 \forall p$, 每个 particle 都服从分布进行传播: $s_{t+1}^p \sim \tilde{f}_{\theta_{b(p,t)}}(s_t^p, a_t)$, 其中 particle 的选择是基于一个 bootstrap $b(p,t) \in \{1, \dots, B\}$, 最终将 P 组 Particles 采集到的结果进行 ensemble。

Algorithm 1 Our model-based MPC algorithm ‘*PETS*’:

- 1: Initialize data \mathbb{D} with a random controller for one trial.
 - 2: **for** Trial $k = 1$ to K **do**
 - 3: Train a *PE* dynamics model \tilde{f} given \mathbb{D} .
 - 4: **for** Time $t = 0$ to TaskHorizon **do**
 - 5: **for** Actions sampled $\mathbf{a}_{t:t+T} \sim \text{CEM}(\cdot)$, 1 to NSamples **do**
 - 6: Propagate state particles \mathbf{s}_τ^p using *TS* and $\tilde{f} | \{\mathbb{D}, \mathbf{a}_{t:t+T}\}$.
 - 7: Evaluate actions as $\sum_{\tau=t}^{t+T} \frac{1}{P} \sum_{p=1}^P r(\mathbf{s}_\tau^p, \mathbf{a}_\tau)$
 - 8: Update $\text{CEM}(\cdot)$ distribution.
 - 9: Execute first action \mathbf{a}_t^* (only) from optimal actions $\mathbf{a}_{t:t+T}^*$.
 - 10: Record outcome: $\mathbb{D} \leftarrow \mathbb{D} \cup \{\mathbf{s}_t, \mathbf{a}_t^*, \mathbf{s}_{t+1}\}$.
-